

# Experiments in data mining, entity disambiguation and how to think data-structures for designing beautiful algorithms

Ekta Grover

1<sup>st</sup> September, 2013




# Structure of this Talk

- How to think through a data mining problem
- 3 problems in data-mining around disambiguation, Natural language processing & Information Retrieval
- Measuring your Model performance - developing a near-real time performance aggregation platform
- Putting it all together

# The only algorithm you will ever need to know

- 1 Write down the problem
- 2 Think real hard
- 3 Write down the solution



# Thinking through Data-Mining problem

- What is the data and how should you represent it ?
- What are the relationships you are interested in ?
- How to define your Objective Function ?
- How will you know that the score represents the direction you want to capture ?
- How will you measure the performance after you are done ?

# Problem # 1 : Entity Disambiguation

## Problem

**So, How much is your network really worth ?**

## What is Disambiguation ?

Information = Related + Unrelated

Extract significant attributes/features of an entity

Relationships between entities & features

Objective function is a just a way to measure - rank/score the features for an entity wrt to other entities

**Disambiguation is not a deterministic problem**

# Problem # 1 : Entity Disambiguation

Guiding question #1

**What form of disambiguation to use for the entities ?**

## Algorithm & approach

- 1 Preprocessing - lower casing & removing punctuation
- 2 Exploratory Analysis - frequency plots (with NLTK freqdist & Pretty table)
- 3 Tokenization to find "keywords" to disambiguate
- 4 Handling stop words & Internationalization
- 5 Frequent item-set mining with support = #of buckets
- 6 Post processing, frequency plots & visualization

**Extensions:** Handling Typo graphical errors and augmenting this data-set with scraping

# Problem # 1 : Talking Specifics & answering the "Why"

## Frequent item-set mining & concept of support

How do we know the number of significant words in an entity name?

- 1 For every company name - strip the 1<sup>st</sup> word and store it as the key of "local" dictionary, with values as list
- 2 On the local dict created in step 1, do frequent item-set mining with support = # baskets - if a word is significant (and uniquely describes an entity), it should be present in all baskets
- 3 Treat all the values in a local dict as a set - If all elements of a local dict are same - no need for item-set mining
- 4 For every key - find the smallest list of tokens (by length) and search the presence of all tokens iteratively till the support appears in less than # baskets
- 5 Replace the words fetched post support to equal all values in that dictionary
- 6 Fetch all values from all keys in the dict of list - flatten this list
- 7 Regular frequency counts on this flattened list of now disambiguated elements

# Problem # 1 : Limitations of this approach

- ➊ **Abbreviations** : Boston Consulting Group & BCG, State Bank of India & SBI - build a referencing domain specific Word Sense Disambiguation (WSD) database
- ➋ **Differently referenced entities** such as "Innovation labs 247 customer private ltd" and "247 innovation labs" : The referencing key will be different for these
- ➌ **Sub-entities** of a company with punctuation "Ebay/Paypal" - the "key" depends on how we choose to treat the punctuation for {/,,:-}
- ➍ Need a module to handle **inadvertent typographical errors** in entity names, names are proper nouns, yet no referencing from Wordnet database - plug in Problem #2

## Extensions

Cross document and Co-reference resolution, Multi-pass learning, smoothing & annealed learning algorithms





**List of original values in the csv file (37 in total)**

[illegible]

Unique values in the original csv file

```
set(['[24]7 Inc - Innovation Labs', 'Innovation Labs, 24/7 Customer Pvt Ltd', 'Innovation Labs @ [24]7, Inc', '[24]7,Innovation Labs', '[24] 7 innovations labs', '[24]7-inc', '24/7, Inc', '[24]7, Inc', '[24]7 Innovation Labs', '[24]7 Inc', '24/7 Customer', '[24]7 Inc. (iLabs)'])
```

### Post removing punctuation & lowercasing

["247 customer", "247 customer", "247 inc", "247 inc", "247 inc", "247 inc", "247 innovation labs", "247 customer", "247 inc", "247 innovation labs", "247inc", "247 inc", "247innovation",  
on labs", "247 inc", "247 inc", "247 inc", "247 inc", "247 inc", "247 inc", "247 inc", "247 inc", "247 inc innovation labs", "247 inc", "innovation labs 24",  
7 customer pvt ltd", "247 customer", "247 inc", "247 inc", "247 inc", "247 inc", "247 inc", "247 inc", "247 customer", "24 7 innovations labs", "innovation labs 247",  
inc", "247 inc", "247 inc", "247 inc", "247 inc"]

We have 5 dicts for this entity, which is a more extreme case need a multi-pass here/ pre-processing prior to frequent item-set mining

Post tokenization & keying-in to define dicts with "key" as the 1st word in the company name and value as the list of values with the same 1st word(or key)

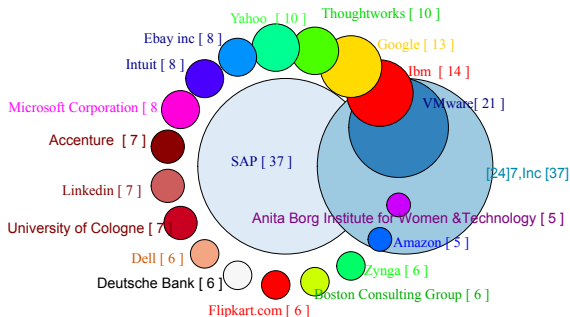
[illegible]

### Post frequent item-set mining without the multipass/pre-processing

[illegible]

<https://github.com/ekta1007/Experiments-in-Data-mining>

# Problem # 1 : Visualization



Top companies by Frequency distribution, post item-set mining

\*I used R(ggplot2 and igraph) to plot the nodes with ColorBrewer palette and Inkspace for "cosmetics"

# Problem # 2 : Custom Distance metric for handling Typographical errors optimized for a QWERTY keyboard

## Usecase

Disambiguating the typographical errors with a contrasting (noun)item

Three types of misspellings<sup>1</sup>

- ① **Typographic errors** - know the correct spelling, but motor coordination slip when typing
- ② **Cognitive errors** - caused by lack of knowledge of the person
- ③ **Phonetic errors**<sup>2</sup> - sound correct, but are orthographically incorrect

Distance measures: Levenshtein, Manhattan, Euclidean, Cosine - inappropriate for distance in our context. Need a **hybrid contextual** approach

---

<sup>1</sup>Karen Kukich - Techniques for automatically correcting words in text, ACM Computing Surveys Volume 24 Issue 4, Dec. 1992 ; Creating a Spellchecker with Solr <http://emmaespina.wordpress.com/2011/01/31/20/>

<sup>2</sup>Soundex algorithm; example : Chebyshev & Tchebycheff

# Problem # 2 : QWERTY distance metric


## Guiding Assumption

Since typographical errors are inadvertently introduced, they should not count towards the overall distance

## Algorithm & approach?

- 1 **Data structures** - QWERTY keyboard as a graph with connected adjacent nodes, neighborhood nodes with appropriate distances
- 2 **Define a typographical error & it's thresholds**- what can be the difference in length of words, how many possible "swaps" can occur in a typo
- 3 **Implement** the custom distance metric - fall back to Levenshtein when non-typographical - Hybrid approach
- 4 **Output** whether the words are identical, or the distance taking typographical errors into account
- 5 **Contrast and blend in** with other fuzzy logic methods, Apache Solr, Spell checker & Auto-suggest

## Problem # 2 : Representation & Codewalk



QWERTY Keyboard representation

► Codewalk from Github

<https://github.com/ekta1007/Custom-Distance-function-for-typos-in-hand-generated-datasets-with-QWERTY-Keyboard/>

# Problem # 3 : Designing your custom job feeds at LinkedIn

## Part 1

Create filtering criteria to **build your custom data-set**

Linkedin REST API's, scraping(Scrapy, BeautifulSoup), raw html with urllib & NLTK

## Part 2

Build **relevance score** of the filtered results

TF-IDF(with cosine similarity), Structured Relevance Models (SRM), Proximity search, classical Information Retrieval & pattern matching algorithms

# Problem # 3 : Part 1

## Algorithm & approach

- ① **Exploit the commonality** in url pattern of the Jobs page, looping constructs for **job\_index** `http://www.linkedin.com/jobs?viewJob=&jobId=job_index&trk=rj_jshp`
- ② Post this **only 4 options** exist per Job page
  - ① The job you're looking for is no longer active
  - ② We can't find the job you're looking for
  - ③ The job passes your filtering criteria (location, skill-specific keywords etc)
  - ④ The job is active, but does not pass your filtering criteria
- ③ For urls in 2.3, **build a simple csv file** with all elements of interest - Company Name, Title of Position, Job Posting Url, Posted on, Functions, Industry, similar Jobs
- ④ Can **pre-filter** by company name - by heuristics, or a quick scoring function for each company name



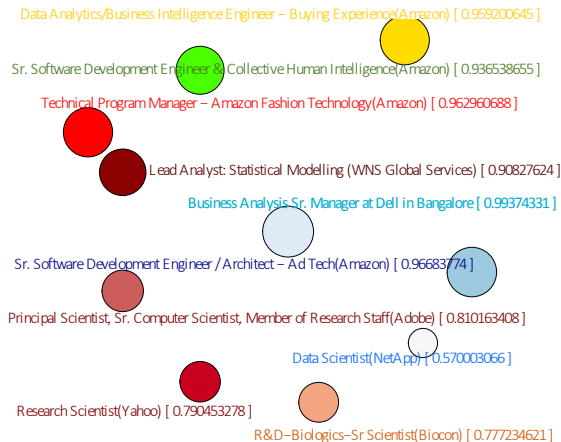
# Problem # 3 : Part 2

## Algorithm & Approach

- 1 A document is a **bag-of-words** - Tokenize, standard preprocessing(punctuation, lower case)
- 2 **Exploratory analysis** with NLTK's Dispersion plot, FreqDist, collocations, common\_contexts, count, concordance
- 3 **Lemmatization & stemming** with NLTK's derivationally related forms with Morphy(Wordnet), Porter, Wordnet\_app
- 4 **TF-IDF** - with stop word list from NLTK & additional weighage for words of your specific interest
- 5 **Inverted index** - for each word in vocabulary, store the doc-id with it's TF-IDF, making up the whole TF-IDF vector
- 6 Given the TF-IDF vector of the document corpus & query, compute the **cosine similarity** of each document with the base query
- 7 **Output top-k** relevant Job listings, with the top-k' words in each job feed
- 8 **Augment** this with a classical structured relevance model, Proximity search or pattern matching algorithm

# Problem # 3 : Representation & Visualization

10



Relevant Job feeds ranked by Cosine Similarity & TF-IDF to the base query

# Problem # 3 : Representation & Codewalk

architect [ 0.006532478 ]  
ad [ 0.005993003 ] amazon [ 0.007088328 ]  
traffic [ 0.006636242 ] distributed [ 0.004794402 ]

## Sr. Software Development Engineer / Architect-Ad Tech(Amazon)

advertising [ 0.015726458 ]  
highly [ 0.008848322 ] increase [ 0.006636242 ]

analyzes [ 0.008369617 ]  
goals [ 0.008369617 ] dell [ 0.02380102 ]

## Business Analysis Sr. Manager(Dell)

internal [ 0.006277213 ] procurement [ 0.023016447 ] strategy [ 0.010412946 ]

develops [ 0.010462021 ] reports [ 0.008369617 ] planning [ 0.006277213 ]

discovering [ 0.013055691 ]  
enable [ 0.009432181 ] nextorbit [ 0.047870868 ]  
space [ 0.00928175 ] energy [ 0.015740528 ]

## Sr. Energy Data Scientist(NextOrbit Inc)

scientist [ 0.01100065 ]  
trends [ 0.00928175 ] understanding [ 0.011015618 ]

Top ranking words in the respective job-feeds<sup>3</sup>

► Codewalk from Github

<https://github.com/ekta1007/Creating-Custom-job-feeds-for-Linkedin>

<sup>3</sup> Note that some of the words are relatively uninformative like "understanding", "Amazon" - could use any amongst Proximity search, Phase search, Positional index, Next word index

# Problem # 4 : Designing a near-real time Performance Aggregation Platform

## Concept

Once you roll your competing candidate models into production, they should self-decide based on performance

Competing Goals can be amongst

- 1 **Minimizing entropy** of the complete system (more centered distribution)
- 2 **Optimizing precision & recall**<sup>4</sup> - tradeoff between type I & type II
- 3 **Accounting for trends** - 1<sup>st</sup> order (velocity) and 2<sup>nd</sup> (acceleration)
- 4 **Minimize bias** - Low mean square error, moving towards a more global representation of a problem

---

<sup>4</sup> precision - fraction of retrieved instances that are relevant, recall - fraction of relevant instances that are retrieved

# Problem # 4 : Defining the aggregated proxy for Performance

## Algorithm & approach

- 1 **Apriori:** Initialize the candidate population, all else equal
- 2 **Define model** equations, which generate the scores, per model
- 3 **Define the metrics** for weighted performance and what is considered a "success event"
- 4 **Define threshold** conditions and measure for success for each metric, wherever applicable
- 5 **Scale the aggregated performance** to have the next period candidate population
- 6 Roll out the new candidate population to next period, per model (Probabilistic model)

Scores are **centrality scores**(Welford's Algorithm) - not to be over-enthusiastic(pessimistic) for state-outcomes and weights for **recency of performance**, over long time learning cycles

# Problem # 4 : Designing the performance benchmark

Time period bench-marked against	Model #	Proxy for bias	Type I & Type II errors		Measure of entropy		Velocity Metric		Acceleration Metric	
		Mean Square Error	Precision	Recall	Mean Score	Mean Variance Score	Velocity by Metric 1	Velocity by Metric 2	Acceleration by Metric 1	Acceleration by Metric 2
t1	M1									
	M2									
	M3									
	M4									
	M5									
t2	M1									
	M2									
	M3									
	M4									
	M5									

# Problem # 4 : Representation & Codewalk



```
Prototype of the Real Time Batch Aggregator
.....
#####Time 1 #####
Model# | Time | #Candidates | Mean_Sq_| Mean | Variance| Precision| Accuracy| Velocity(V) |Velocity(C) |Acceleration(V) |Acceleration(C)
          Error (Score) (Score)          (visit)      (chats)      (Visit)      (chats)
-----
Model# 1| 1 | 14000 | 68.431776 | 0.501543 | 0.082948 | 0.148243 | 0.611786 | 0.075000 | 0.750000 | 0.000000 | 0.000000
Model# 2| 1 | 14000 | 68.477386 | 0.501891 | 0.083045 | 0.089988 | 0.602214 | 0.045000 | 0.281250 | 0.000000 | 0.000000
Model# 3| 1 | 14000 | 68.481001 | 0.500558 | 0.084417 | 0.042317 | 0.404500 | 0.030000 | 0.300000 | 0.000000 | 0.000000
Model# 4| 1 | 14000 | 68.762492 | 0.503686 | 0.084035 | 0.099326 | 0.398286 | 0.075000 | 3.750000 | 0.000000 | 0.000000
Model# 5| 1 | 14000 | 68.190888 | 0.499248 | 0.082894 | 0.055354 | 0.550000 | 0.030000 | 0.500000 | 0.000000 | 0.000000
#####PERFORMANCE AGGREGATED METRIC #####
Model# | Time | #Candidates | Performance Metric
          (Aggregated & Normalized)
-----
Model# 1| 1 | 14000 | 4.10874748831
Model# 2| 1 | 14000 | 2.93421049334
Model# 3| 1 | 14000 | 1.83716273545
Model# 4| 1 | 14000 | 5.28702053689
Model# 5| 1 | 14000 | 0.89674366591
.....
delta t - Waiting to Aggregate,Process & Roll the next period decisions
.....
```

## Aggregated performance in Period 1

► Codewalk from Github

<https://github.com/ekta1007/Performance-aggregation-platform—learning-in-near-real-time>

# Problem # 4 : Representation cont..



Aggregated performance in Period 1 & 2 respectively



# Resources

- ▶ [Link to the problems I talked today\(I keep updating this\)](#)
- ▶ [Understanding data structures in Python](#)
- ▶ [A Programmer's Guide to Data Mining - Ron Zacharski](#)
- ▶ [Creating a Spell Checker with Solr](#)
- ▶ [An Intuitive example to TF-IDF, Iraq War logs - Jonathan Stray](#)
- ▶ [An Introduction to Information Retrieval - Christopher D. Manning, Prabhakar Raghavan, Hinrich Schutze](#)
- ▶ [Natural Language Processing with Python - Steven Bird, Ewan Klein, Edward Loper](#)
- ▶ [Stemming & Lemmatization](#)
- ▶ [Programming Collective Intelligence: Building Smart Web 2.0 Applications - Toby Segaran](#)
- ▶ [Mining the Social Web: Analyzing Data from Facebook, Twitter, LinkedIn - Matthew A. Russell](#)



Questions  
are  
guaranteed in  
life;  
Answers  
aren't.